

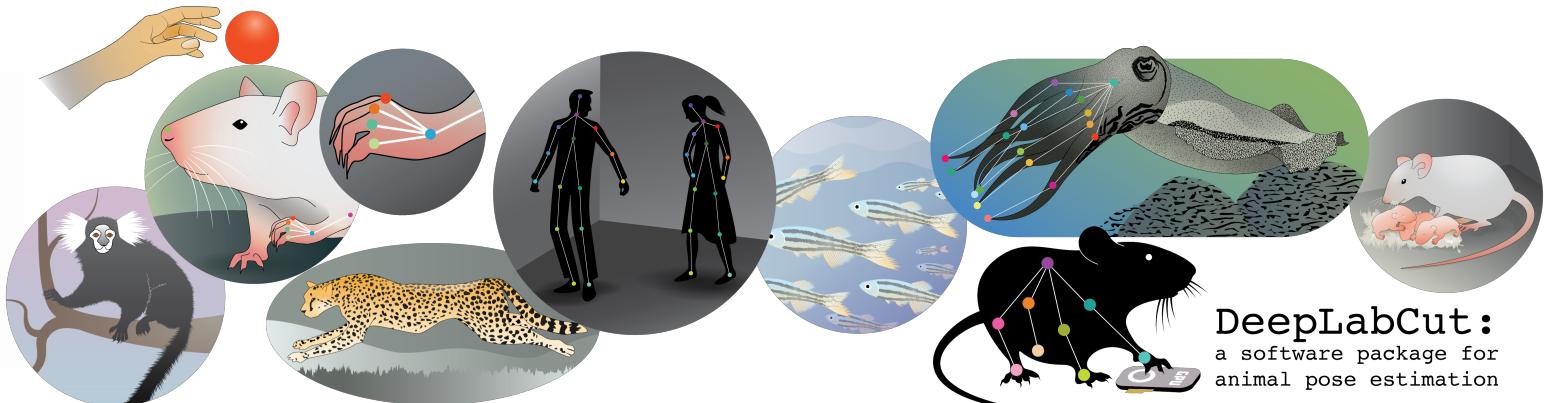
DeepLabCut workshop

Session 2: training & network selection

16th and 17th of January 2020

Rowland Institute, Harvard University

Mackenzie Mathis
Alexander Mathis
Jessy Lauer



Chan
Zuckerberg
Initiative 

HARVARD
UNIVERSITY



“Software 2.0” – integration of annotation, training and inference

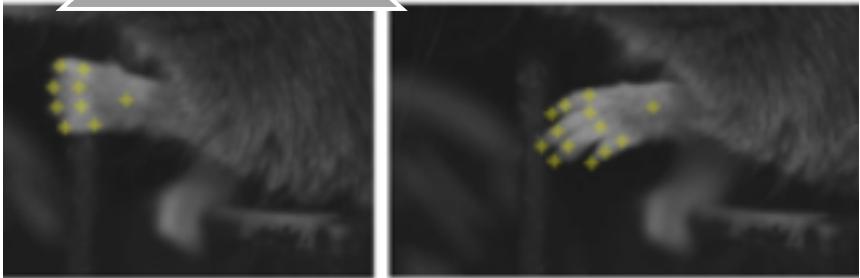
Train DNN

Create a project,
extract frames, +
GUIs to label your data

Select + Train your
deep neural network

Evaluate network
performance
(active learning + GUIs
if improvement needed)

Run inference on
new videos,
create labeled videos,
+ plot your results!



refine?



What network & augmentation method?

Network architecture:

- Network backbones: MobileNets V2, *ResNets*
- Scale (level of deconvolution)
- Differences: **Inference & training speed, resolution**

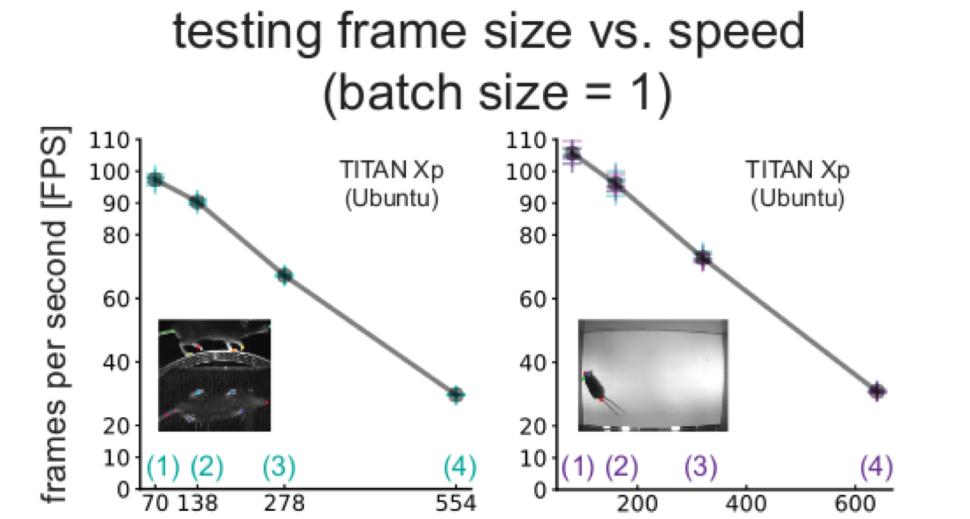
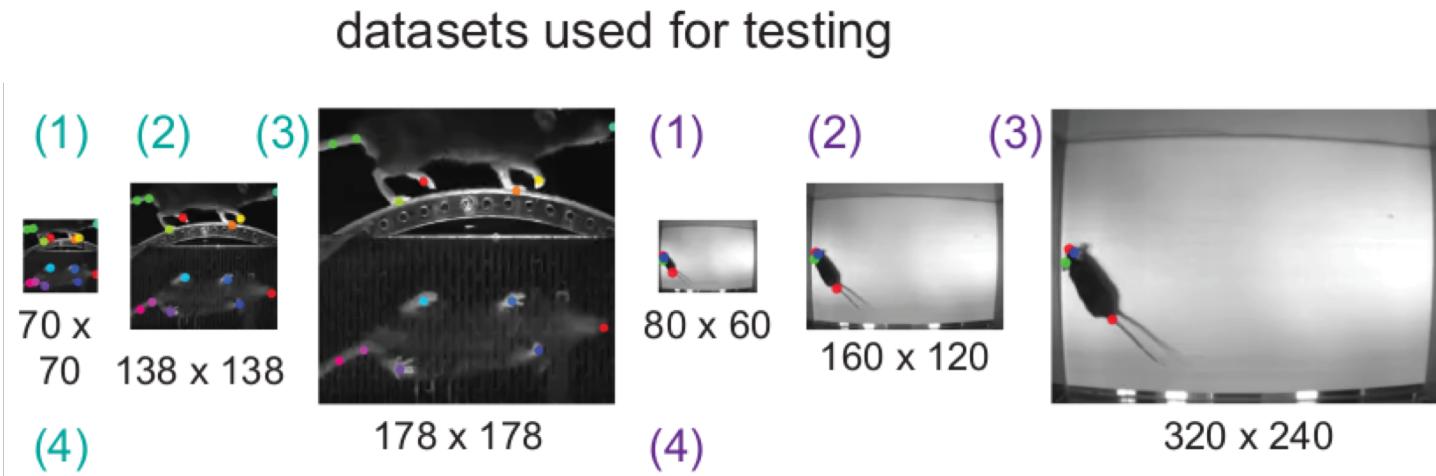
Pre-training (Weights):

- *ImageNet pretrained (the generalist)*
- MPII pose pretrained
- DeepLabCut model zoo....

Training method:

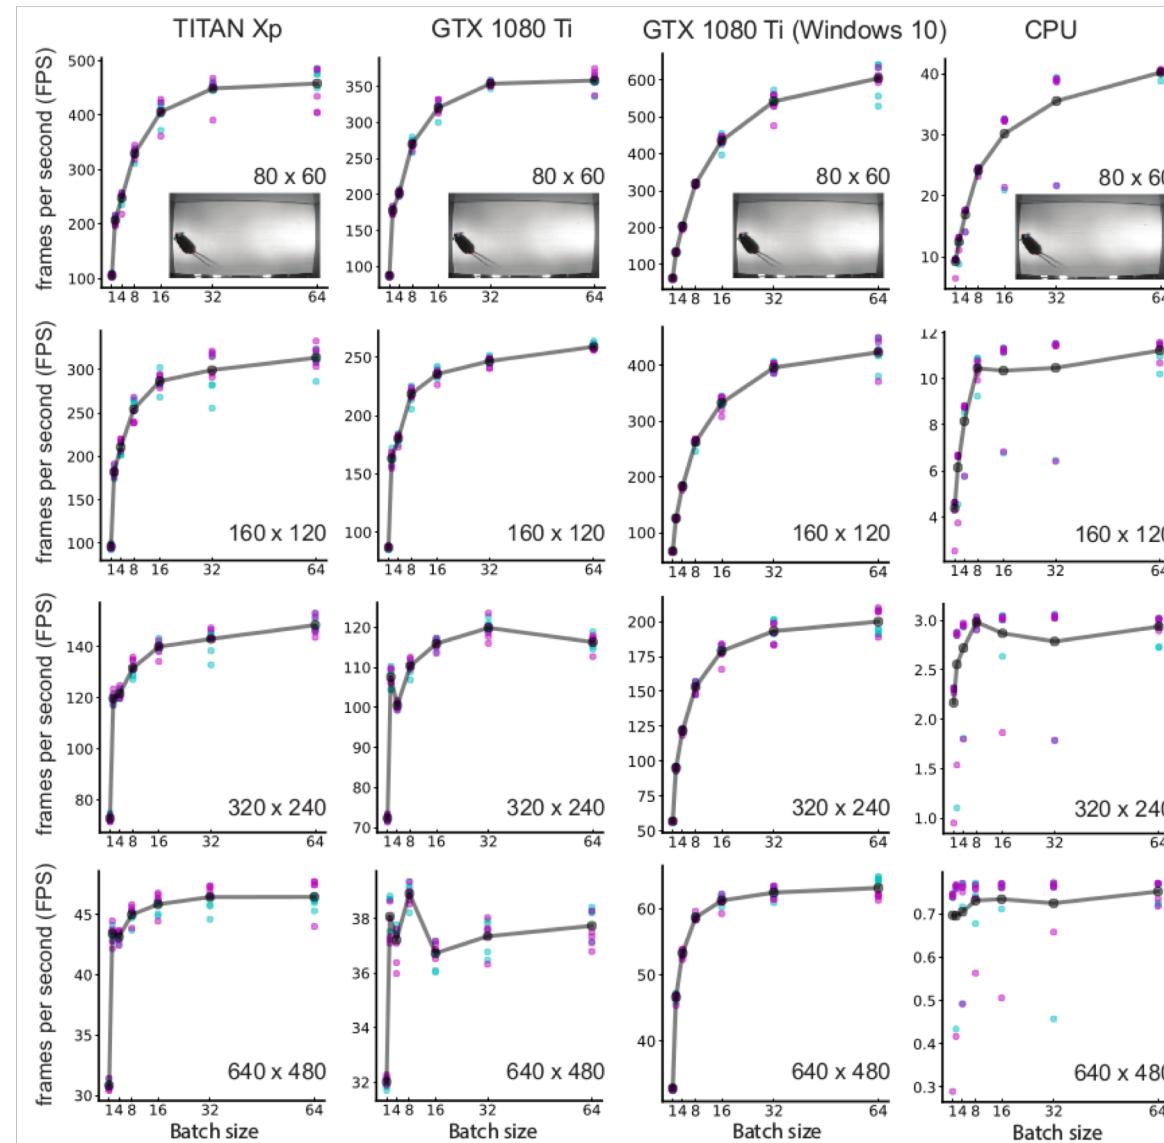
- Learning rate & optimizer (Adam/SGD, batch_size, ...)
- Augmentation transformations (rotation, scaling, ...)

Inference speed strongly depends on image size

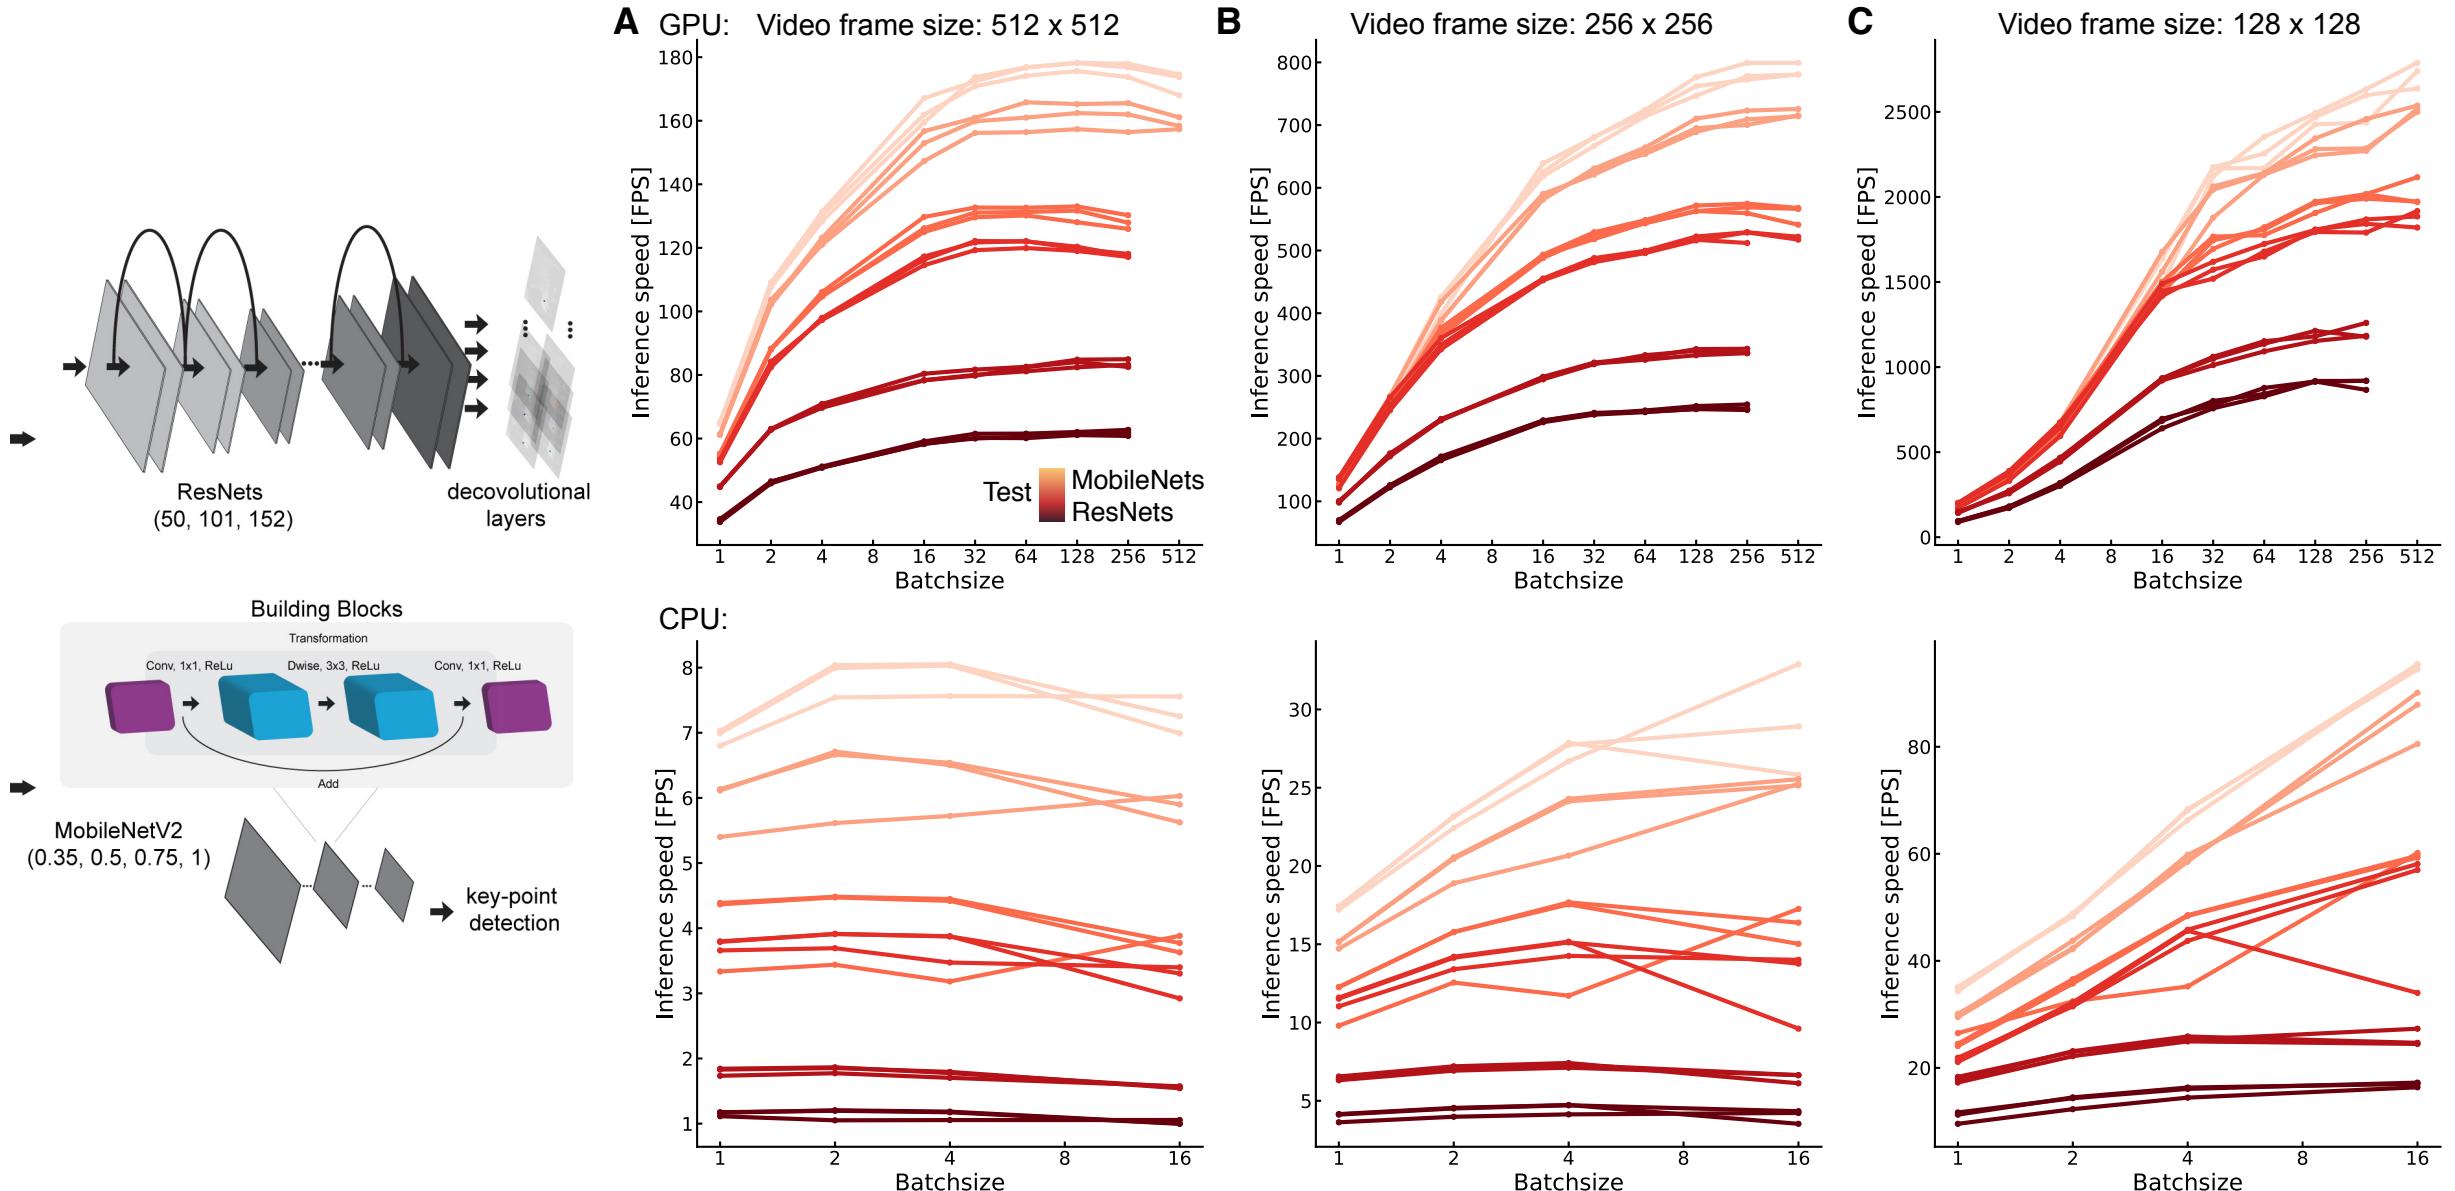


Inference speed depends on hardware!

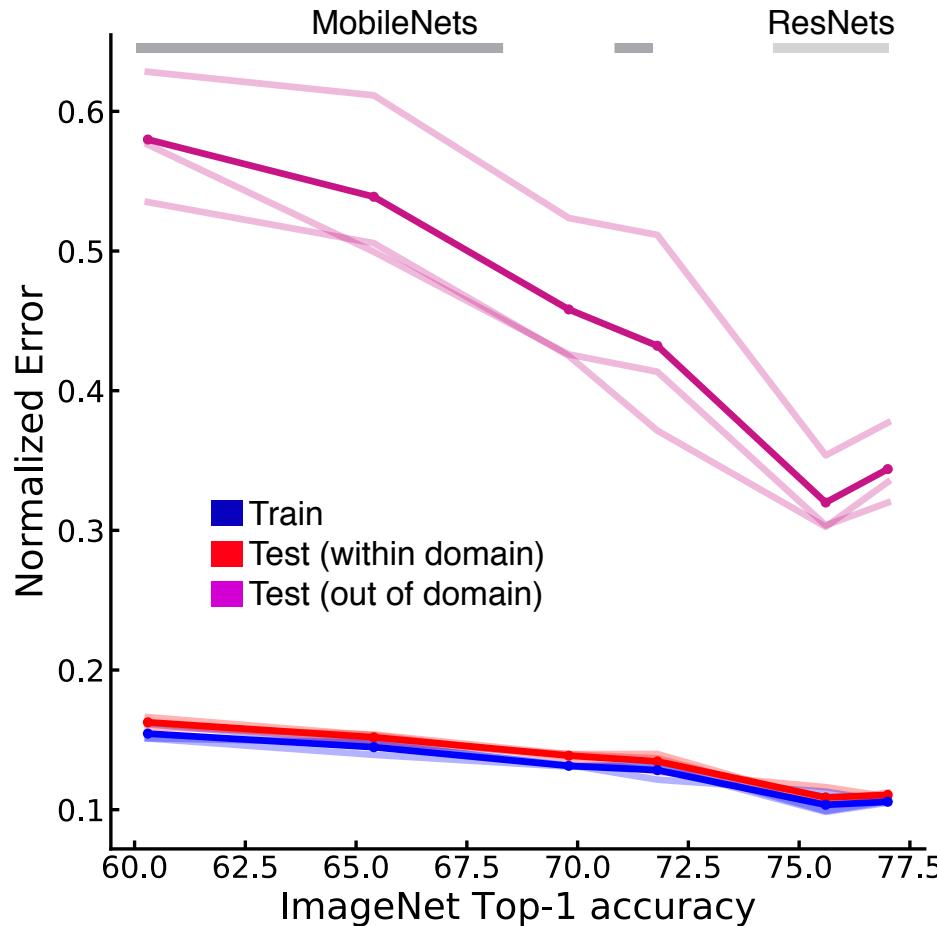
- Get a GPU!



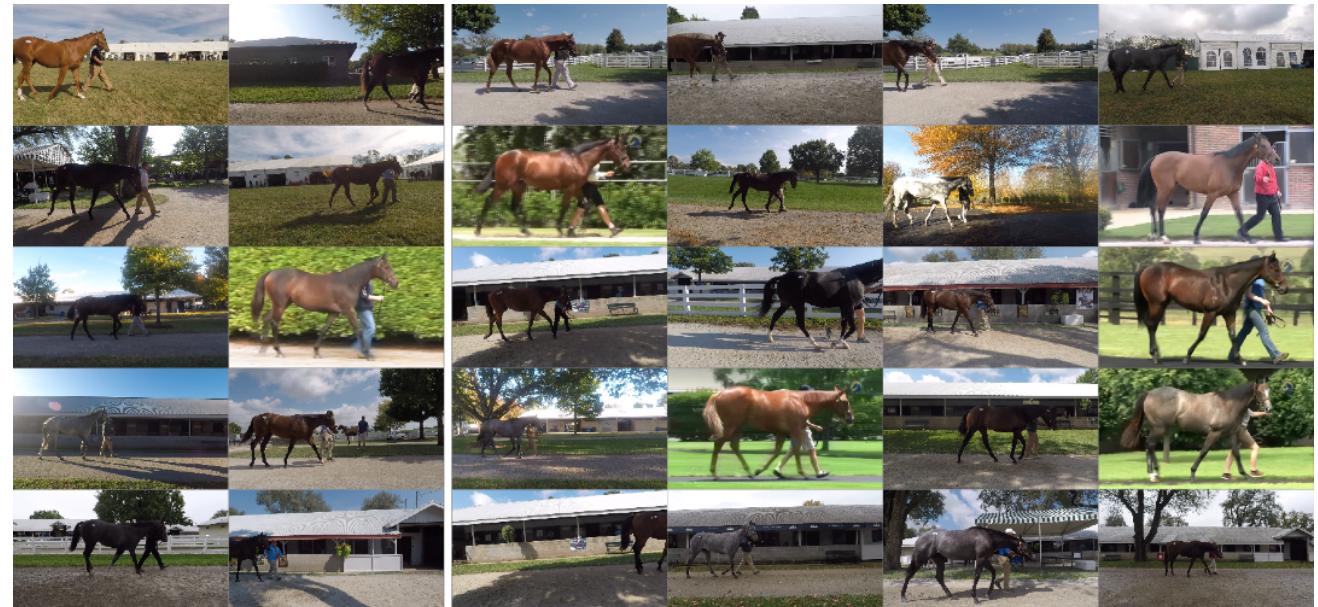
MobileNets are faster than ResNets



ImageNet accuracy correlates with performance



Within domain: horses as contained in data set
Out of domain: different horses & background

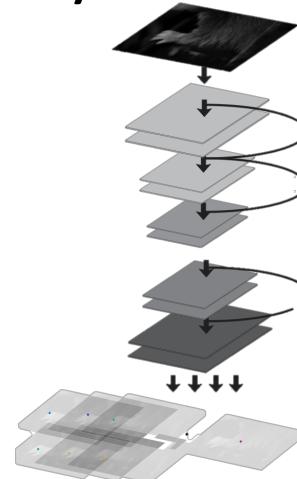


How to pick the network type?

- **Performance vs. speed trade-off.**
Higher robustness requirements: ResNet 50 (or 101),
High speed requirements: MobileNetsV2 0.35
- Try to down sample/crop videos (the smaller videos, the faster irrespective of model)
- How to choose in DLC?
 - By default uses “default_net_type” from config file of project!
 - Otherwise during **training set creation** you can pass your choice, e.g.:
`deeplabcut.create_training_dataset(config, net_type='mobilenet_v2_1.0')`
Pick from *resnet_50, resnet_101, resnet_152, mobilenet_v2_1.0, mobilenet_v2_0.75, mobilenet_v2_0.5, mobilenet_v2_0.35*
Soon: EfficientNet
- Alternative: change **net_type** in **config_pose.yaml** (for both train & test folders!)

What type of pre-training (weights)?

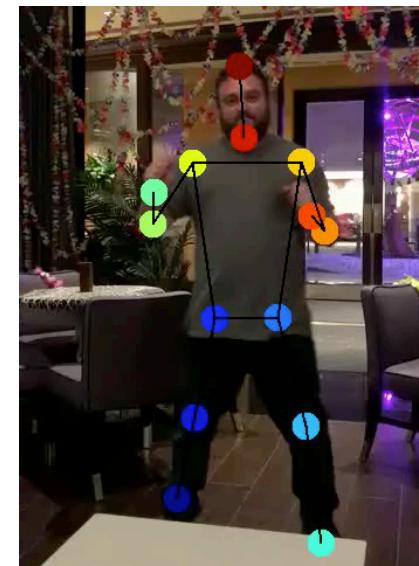
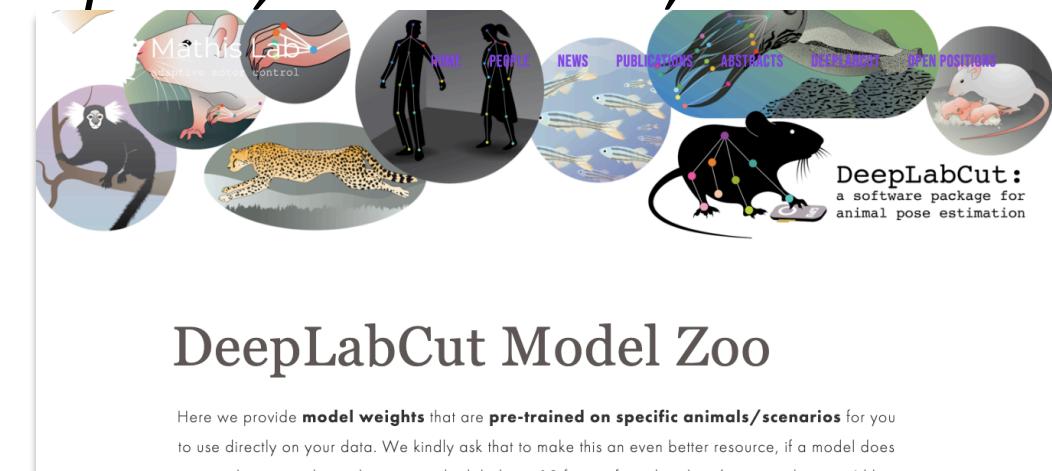
- Why **pre-training?** (*faster training, less data required, more robust;* see Mathis et al. 2019 arxiv)
- Currently: “generic” ImageNet pre-training
- **Use weights from previous training/others!**
- Soon: *pre-trained models for typical lab task*
- Already available: MPII pose pretrained



Never re-trained >

DeepLabCut includes the ability to use
ResNet-101 pre-trained on MPII
26K images, 40K humans

Code: `deeplabcut.create_pretrained_human_project(...)`



For loading different weights:
Set path in
config_pose.yaml file!

What type of training/augmentation?

- Batch size (the larger the better; but large memory demand)
- Learning schedule Adam, SGD; rates are good
- **Augmentation**

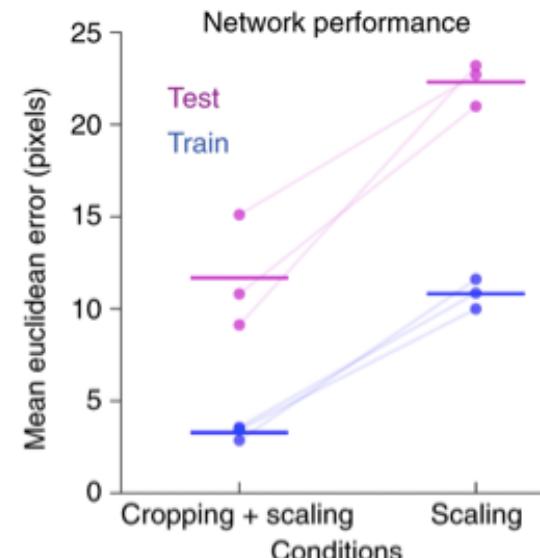
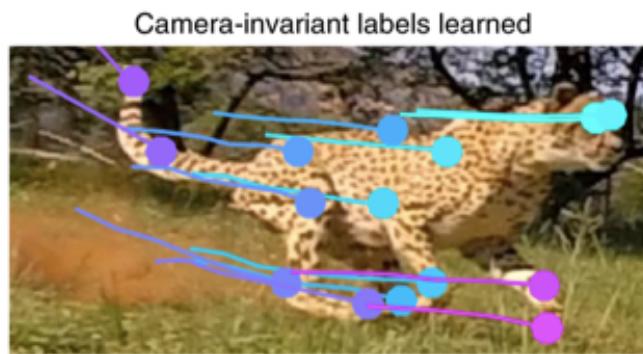
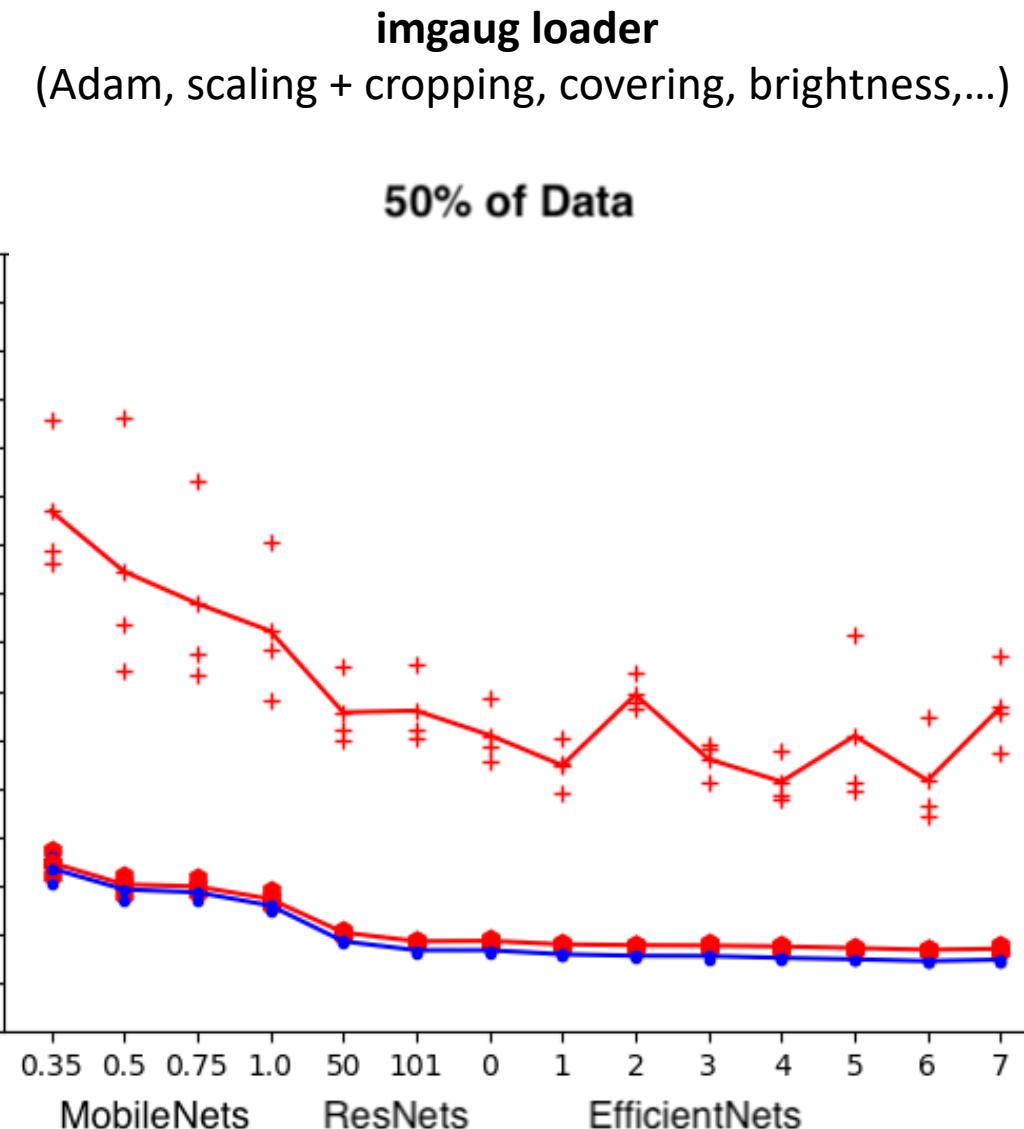
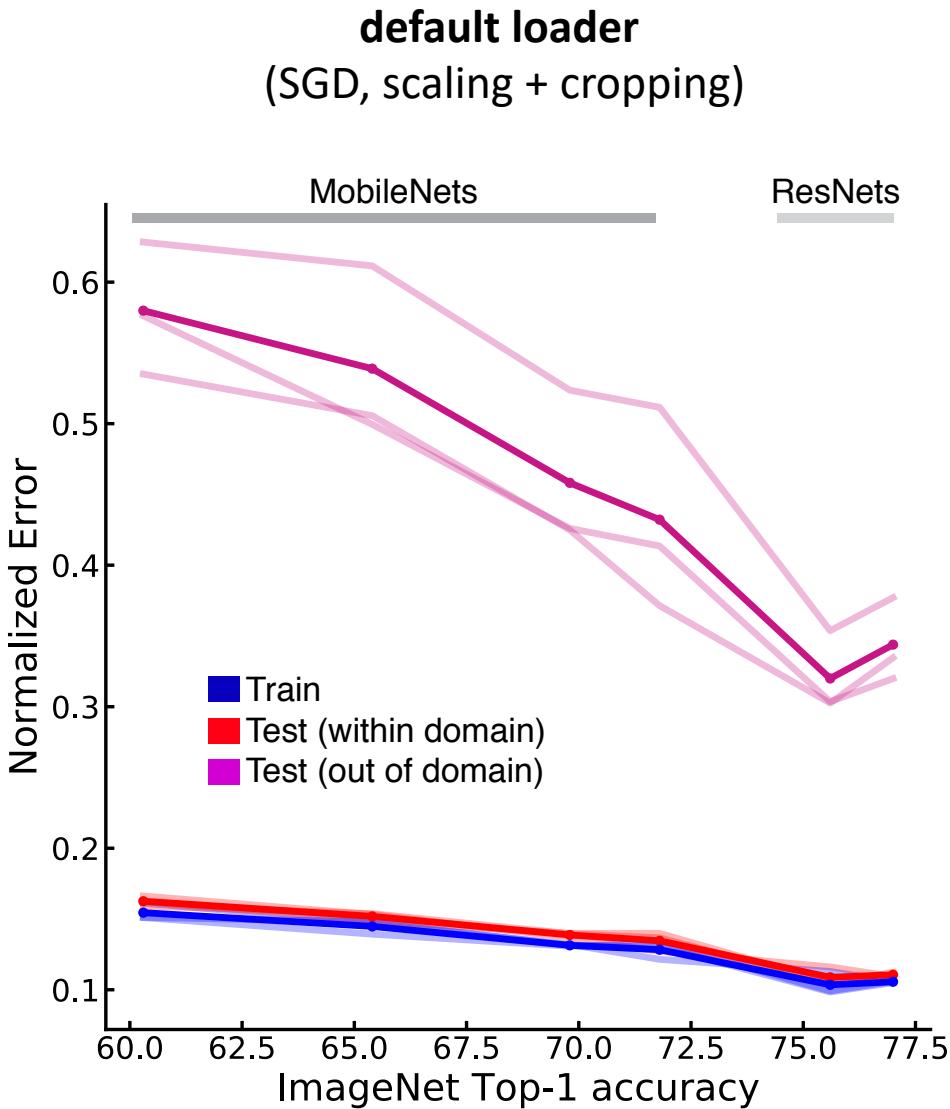


	Image	Keypoints
Original Input		
Gauss. Noise + Contrast + Sharpen		
Affine		
Crop + Pad		
Fliplr + Perspective		

Augmentation examples from Imgaug

The impact of augmentation methods



How to set augmentation methods in DLC?

By default uses “`default_augmenter`” in config file of project!

Alternatively, during **training set creation** you can pass your choice,
e.g.:

```
deeplabcut.create_training_dataset(config, augmenter_type='imgaug')
```

Pick from *default*, *imgaug*, *tensorpack* & *deterministic*

You can also change **dataset_type** in **config_pose.yaml** as well as other variables of interest (i.e. `global_scale`, `multi_step`, ...).

See BOX 2 in Nath*, Mathis* et al for details

Box 2 | Parameters of interest in the network configuration file, *pose_cfg.yaml*

Please note, there are more parameters that typically never need to be adjusted; they can be found in the default *pose_cfg.yaml* file at https://github.com/AlexEMG/DeepLabCut/blob/master/deeplabcut/pose_cfg.yaml.

- **display_iters**: An integer value representing the period with which the loss is displayed (and stored in *log.csv*).
- **save_iters**: An integer value representing the period with which the checkpoints (weights of the network) are saved. Each snapshot has >90 MB, so not too many should be stored.
- **init_weights**: The weights used for training. Default: <DeepLabCut_path>/Pose_Estimation_Tensorflow/pretrained/resnet_v1_50.ckpt. For ResNet-50 or 101, -- this will be automatically created. The weights can also be changed to restart from a particular snapshot if training is interrupted, e.g., <full path>-snapshot-5000 (with no file-type ending added). This would re-start training from the loaded weights (i.e., after 5,000 training iterations, the counter starts from 0).
- **multi_step**: These are the learning rates and number of training iterations to perform at the specified rate. If the users want to stop before 1 million, they can delete a row and/or change the last value to be the desired stop point.
- **max_input_size**: All images larger with size width × height > max_input_size*max_input_size are not used in training. The default is 1500 to prevent crashing with an out-of-memory exception for very large images. This will depend on your GPU memory capacity. However, we suggest reducing the pixel size as much as possible; see Mathis and Warren²⁷.

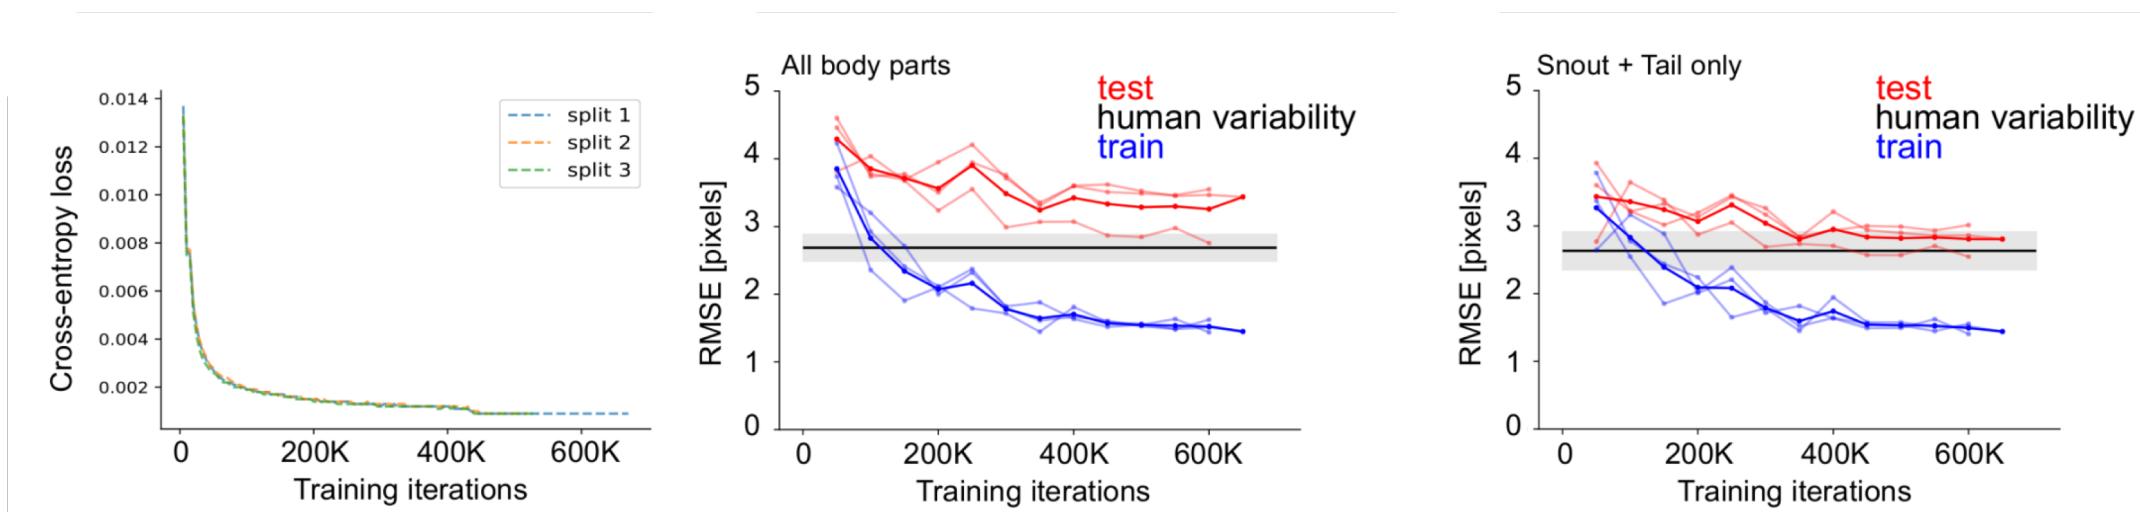
The following parameters allow one to change the resolution:

- **global_scale**: All images in the dataset will be rescaled by the following scaling factor to be processed by the convolutional neural network. You can select the optimal scale by cross-validation (see discussion in Mathis et al.¹²). Default is 0.8.
- **pos_dist_thresh**: All locations within this distance threshold (measured in pixels) are considered positive training samples for detection (see discussion in Mathis et al.¹²). Default is 17.

The following parameters modulate the data augmentation. During training, each image will be randomly rescaled within the range [scale_jitter_lo, scale_jitter_up] to augment training:

- **scale_jitter_lo**: 0.5 (default).
- **scale_jitter_up**: 1.5 (default).
- **mirror**: If the training dataset is symmetric around the vertical axis, this Boolean variable allows random augmentation. Default is False.
- **cropping**: Allows automatic cropping of images during training. Default is True.
- **cropratio**: Fraction of training samples that are cropped. Default is 40%.
- **minsize**, **leftwidth**, **rightwidth**, **bottomheight**, **topheight**: These define dimensions and limits for auto-cropping.

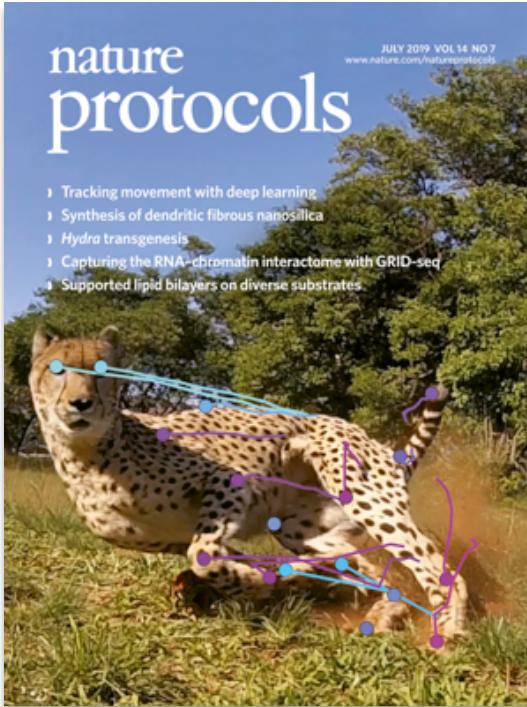
Training considerations



Mathis et al. Nature Neurosci, 2018

- Cross entropy loss minimized (measures difference in probabilities between epsilon disk around ground-truth and predicted scoremap)
- Training speed (mostly) depends on frame size! (and GPU, CPU type)
- Train until the “loss flattens”
- Store as many **snapshots** as necessary (intermediate weight steps; consider early stopping)

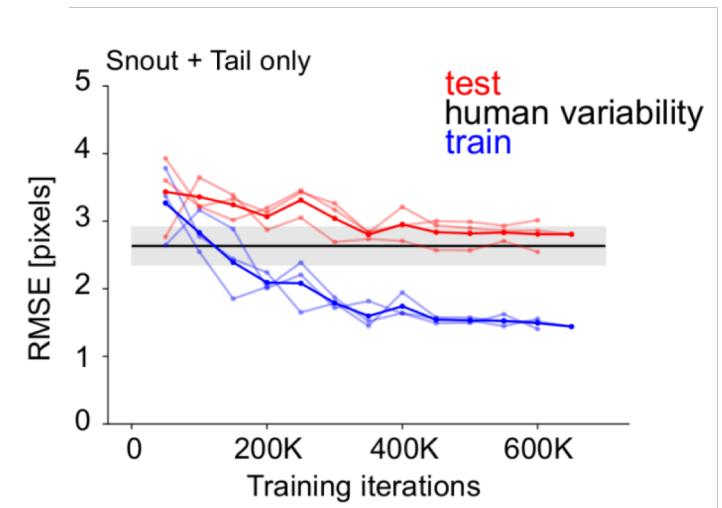
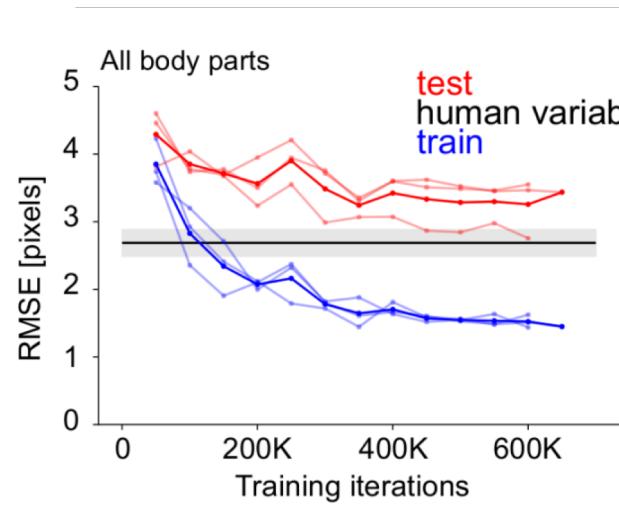
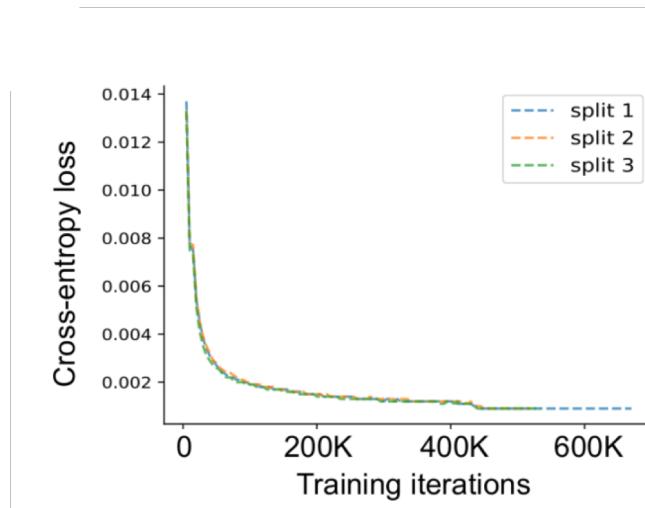
Tutorial: create your own training dataset + start training



- » Tracking movement with deep learning
- » Synthesis of dendritic fibrous nanosilica
- » Hydra transgenesis
- » Capturing the RNA–chromatin interactome with GRID-seq
- » Supported lipid bilayers on diverse substrates.

The image shows a thumbnail for a research article. At the top left is the 'nature protocols' logo. To the right is the word 'PROTOCOL'. Below that is the DOI link: <https://doi.org/10.1038/s41596-019-0176-0>. The main title of the article is 'Using DeepLabCut for 3D markerless pose estimation across species and behaviors'. Below the title is the author list: Tanmay Nath^{1,5}, Alexander Mathis^{1,2,5}, An Chi Chen³, Amir Patel³, Matthias Bethge⁴ and Mackenzie Weygandt Mathis^{1*}.

Evaluation: Compare test RMSE and ideally your labeling error!



Evaluation >> look at test images!

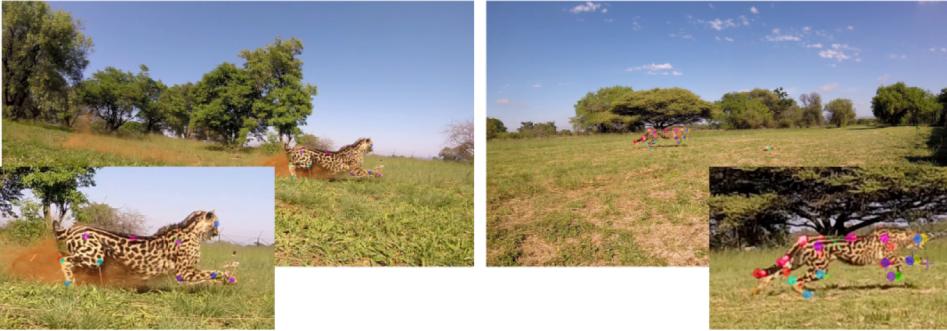
a Example of the terminal output

```
deeplabcut.evaluate_network('<path of the proj. config file>', shuffle = [1], plotting=True)
Assessing accuracy of shuffle # 1 with 99 % training fraction.
Found the following training snapshots: [(200000, 0)]
You can choose among those for analysis of train/test performance.
Results for 200000 training iterations: 99.1 train error: 4.81 pixels. Test error: 7.02 pixels.
With cutoff of 0.1 train error: 3.3 pixels. Test error: 5.2 pixels
```

b Examples of **training images with Human and DeepLabCut labels**



c Examples of **test images with Human and DeepLabCut labels**



d Examples of **test images with Human and DeepLabCut labels**

+ Human applied Label
● DeepLabCut

